

PA01: Processes and Threads

Mudit Kumar (MT25073) | Graduate Systems (CSE638)

1. GitHub Repository

The source code and artifacts are hosted at:

https://github.com/Muditkumar123/MT25073-GRS-PA_1/

2. Implementation Overview

This assignment explores the performance differences between Processes (fork) and Threads (pthread) under three workloads:

- CPU Intensive: Complex math calculations stressing the ALU.
- Memory Intensive: 5MB buffer allocation with stride access to trigger TLB misses.
- I/O Intensive: 128KB repeated writes using 'fsync()' to force physical disk I/O.

3. Part C: Baseline Observations (2 Workers)

Baseline performance metrics for 2 concurrent workers pinned to a single core:

| Program | Mode | Time(s) | CPU(%) | Read(KB) | Write(KB) |
|-----------|------|---------|--------|----------|-----------|
| Program_A | cpu | 6.3 | 101.0 | 2756 | 180 |
| Program_A | mem | 6.4 | 100.0 | 0 | 32 |
| Program_A | io | 14.22 | 26.5 | 12 | 983820 |
| Program_B | cpu | 6.32 | 100.0 | 64 | 64 |
| Program_B | mem | 6.17 | 100.0 | 0 | 28 |
| Program_B | io | 15.22 | 22.0 | 0 | 983860 |

Key Observations:

- I/O Bottleneck: The IO task took significantly longer (~15s) compared to CPU/Mem (~6s), confirming the workload was successfully bound by disk speed.
- Massive Writes: The 'fsync' implementation successfully forced massive writes (~980 MB) to the disk.
- Zero Disk Reads: Despite reading the file back, iostat reported 0 KB Reads. This is because the Linux kernel served the data directly from the Page Cache (RAM).

PA01: Processes and Threads

Mudit Kumar (MT25073) | Graduate Systems (CSE638)

4. Part D: Scalability Analysis

We scaled Program A (2-5 processes) and Program B (2-8 threads) to analyze performance trends.

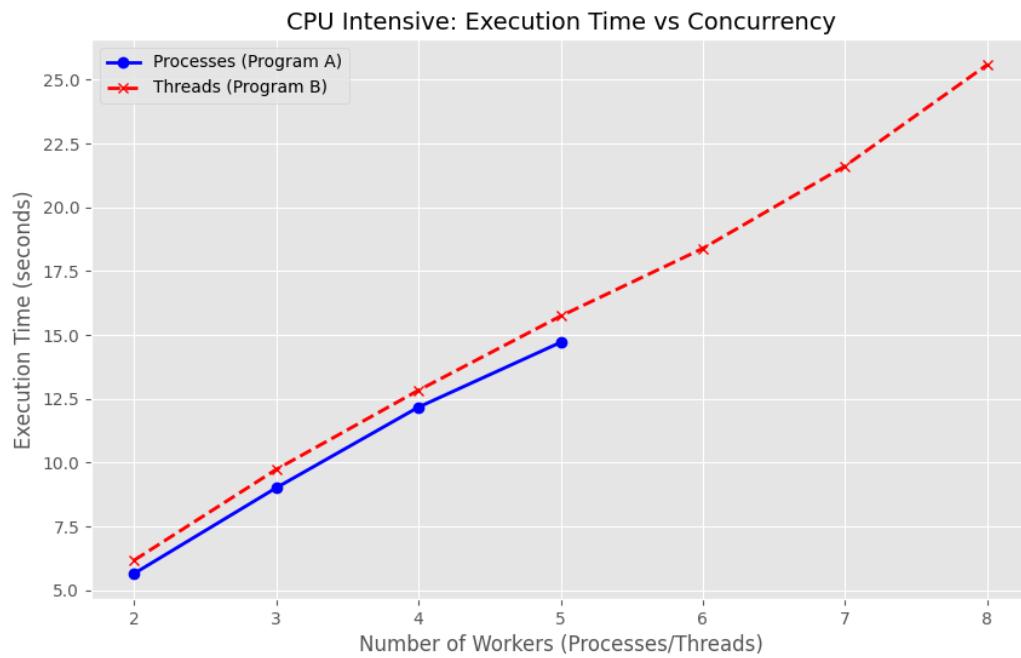


Figure 1: CPU Intensive Scalability

Analysis: Execution time increases linearly. Since workers were pinned to a single core, they executed sequentially (time-slicing).

PA01: Processes and Threads

Mudit Kumar (MT25073) | Graduate Systems (CSE638)

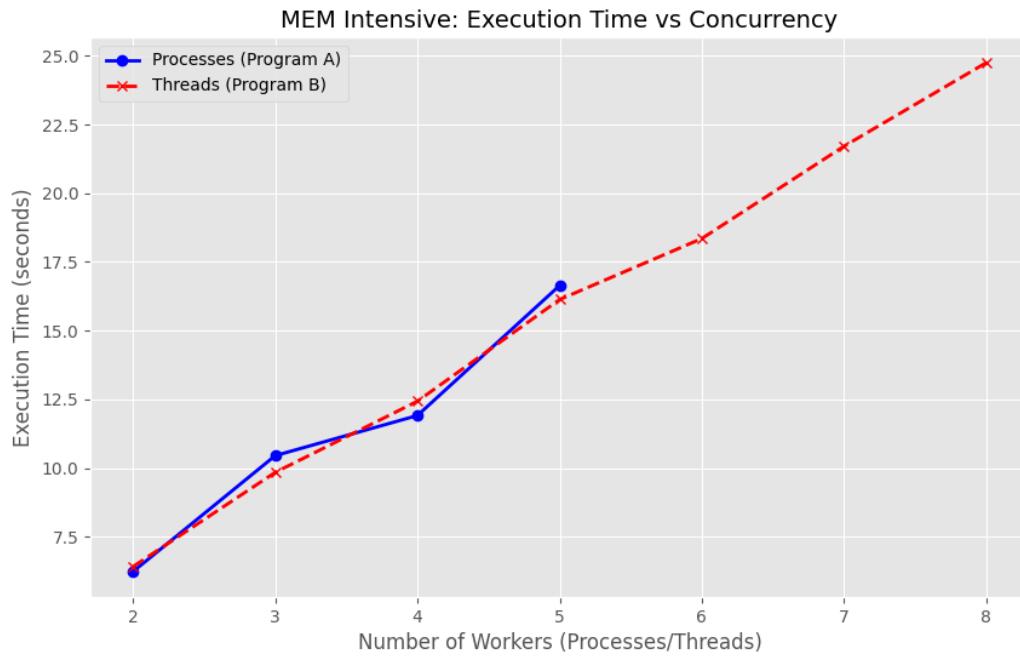


Figure 2: Memory Intensive Scalability

Analysis: Similar linear trend. Threads and Processes showed comparable overhead profiles for memory access.

PA01: Processes and Threads

Mudit Kumar (MT25073) | Graduate Systems (CSE638)

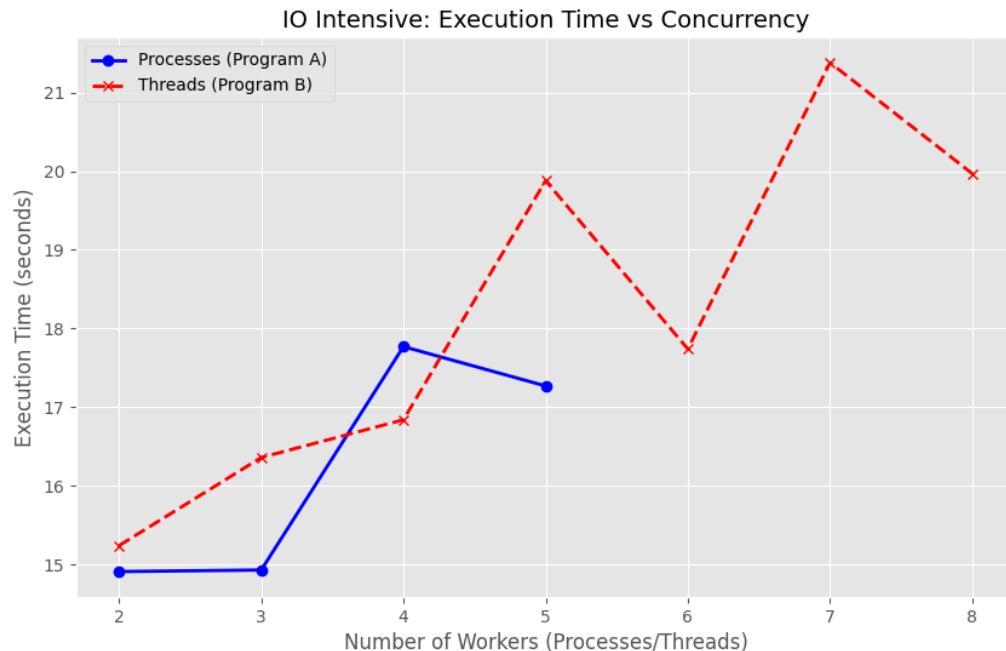


Figure 3: I/O Intensive Scalability

Analysis: The curve is flatter (15s-20s). The disk hit saturation early; adding more threads did not speed up physical disk rotation.

PA01: Processes and Threads

Mudit Kumar (MT25073) | Graduate Systems (CSE638)

5. AI Usage Declaration

I utilized an AI assistant to support the development of this assignment. My usage was as follows:

1. C Code Development: The AI generated the initial code structure and logic for the worker functions. I manually transcribed and verified this code line-by-line to ensure I fully understood the implementation logic.
2. Scripting Support: The AI provided logic for Bash scripts (seq loops, taskset).
3. Debugging: The AI helped diagnose WSL2-specific issues, such as identifying correct iostat columns.
4. Visualization: The AI provided the Python script to generate plots from CSV data.

All final analysis and conclusions are based on the data I collected.